



University of Antwerp
| Faculty of Science

Graph-Based Game Recommendations Using Steam User-Item Interactions

Pablo Deputter

Research Questions

- **How do graph-based recommendation approaches compare to traditional algorithms in improving recommendation accuracy on the Steam platform, and what are the associated computational trade-offs?**
 - *Hypothesis:* Graph-based methods significantly outperform traditional algorithms in recommendation accuracy by capturing complex user-item interactions, albeit with increased computational requirements
- **How does incorporating Personalized PageRank (PPR) enhance recommendation accuracy in game recommendation systems on the Steam dataset?**
 - *Hypothesis:* Personalized PageRank leverages the underlying graph structure of user-item interactions to provide more relevant and accurate recommendations, thereby increasing user satisfaction
- **(How does multi-hop message passing in Graph Neural Networks (GNNs) influence the accuracy and effectiveness of game recommendations on the Steam platform?)**
 - *Hypothesis:* Multi-hop message passing enables GNNs to capture higher-order relationships within the user-item graph, leading to improved recommendation accuracy and better handling of complex interaction patterns)

Stakeholders

- Primary Stakeholder: **Users or Gamers**
 - Enhance user experience by providing accurate and personalized game recommendations
- Secondary Stakeholder: **Game Developers/Publishers**
 - Support publishers by connecting their games with the right audience

Metrics

- Focus on accuracy!
- **NDCG@20**
- **Recall@20**
- **Execution time (s):**
 - Model training + inference (generating recommendations)
 - Are complex models worth it?

Offline Evaluation Setup


- Resource constraints (limited computational power and time)
- Need for efficient evaluation!
- **Sample-Based Evaluation:**
 - Utilize a representative subset of data
 - Ensure reliability despite reduced data size
- **Advantages:**
 - Quick testing of different algorithms
 - Reduced computational resources requirement
 - Ability to explore broader parameter space
- **Disadvantages:**
 - Potential risk of missing patterns present in full dataset
 - Possible suboptimal parameter choices

Offline Evaluation Setup

- Choose **sample size** (number of users)
- Maintain user activity and item popularity distributions
- Data is split in training, validation and test set
 - No overlap across sets to prevent data leakage
- Cross-validation setup
 - Ensure consistent splits for all algorithms
 - Multiple random seeds to ensure robustness
- Fair comparison between algorithms and hyperparameter tuning
 - Same sampled data and splits
 - Only optimization path differs

Baselines

- Offline Evaluation:
 - 3-Fold validation on full training set
 - Metrics computed on validation set
- Online Evaluation:
 - Models trained on full training set
 - Evaluated on new users (Cold-Start)
- Random (0.002)**
- Popularity (0.208)**
 - Surprisingly effective
- User-KNN (0.208)**
 - Cosine similarity between users
 - Fallback to popularity for Cold-Start
- Item-KNN (0.320)**



Personalized PageRank - Introduction





- Originally designed by Google to rank web pages
 - High number of incoming links signifies importance
 - Outgoing links distribute this importance
- **PPR** adapts this by considering user's interaction history → Personalized recommendations
- Natural fit for data
 - Sparse interaction data
 - Cold-start problems
 - Playtime integration
 - Balances popularity with personalization

Personalized PageRank – Parameters

- **α (Alpha):**
 - Damping factor, probability of continuing a random walk
- **num_iterations:**
 - Maximum iterations for the power method
- **popularity_weight:**
 - Balances item popularity in recommendations
- **process_playtime:**
 - How playtime data is processed

Personalized PageRank – Data Processing

- **Log transformation** to reduce outlier impact
 - Normalization using MinMaxScaler
- **Relative scaling** based on user's average playtime



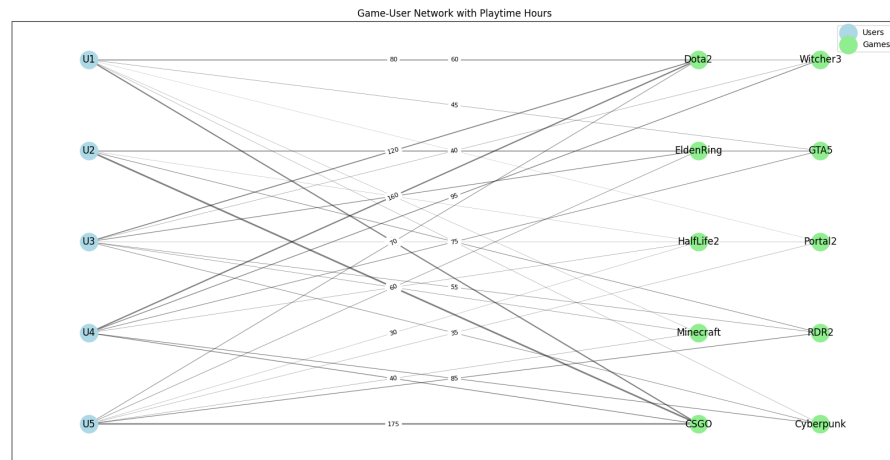
Personalized PageRank – Example

- **Users:** U1, U2, U3, U4 and U5
- **Games:** CSGO, Dota 2, Elden Ring, Minecraft, GTA V, Red Dead Redemption 2, The Witcher 3, Portal 2, Half-Life 2 and Cyberpunk 2077

$$W = \begin{bmatrix} 150 & 80 & 0 & 30 & 45 & 0 & 60 & 20 & 0 & 35 \\ 200 & 0 & 40 & 0 & 90 & 70 & 0 & 0 & 25 & 0 \\ 0 & 120 & 85 & 50 & 0 & 55 & 40 & 30 & 0 & 65 \\ 90 & 160 & 0 & 0 & 75 & 0 & 95 & 0 & 40 & 80 \\ 175 & 70 & 60 & 40 & 0 & 85 & 0 & 35 & 30 & 0 \end{bmatrix}$$



$$W' = \begin{bmatrix} 0.87 & 0.60 & 0 & 0.17 & 0.35 & 0 & 0.47 & 0 & 0 & 0.24 \\ 1.00 & 0 & 0.30 & 0 & 0.65 & 0.54 & 0 & 0 & 0.10 & 0 \\ 0 & 0.78 & 0.63 & 0.40 & 0 & 0.44 & 0.30 & 0.17 & 0 & 0.51 \\ 0.65 & 0.90 & 0 & 0 & 0.57 & 0 & 0.67 & 0 & 0.30 & 0.60 \\ 0.94 & 0.54 & 0.47 & 0.30 & 0 & 0.63 & 0 & 0.24 & 0.17 & 0 \end{bmatrix}$$



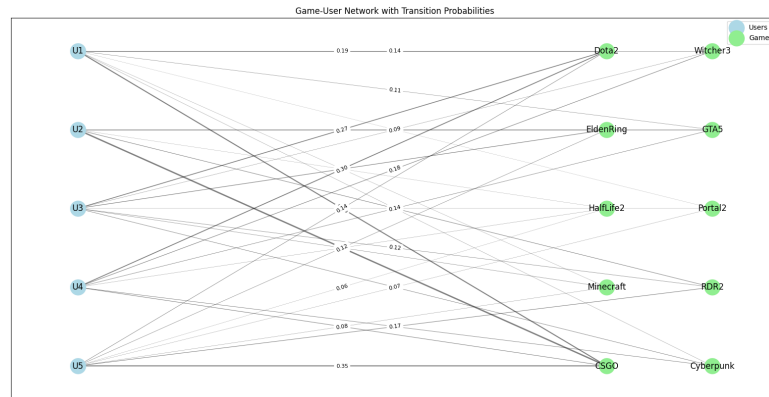
Personalized PageRank - Components

- Interaction (User-Item) Matrix:**

- Sparse matrix with normalized playtime
- Rows sum to 1 representing transition probability from user \rightarrow item

- Item-User Matrix:**

- Transposed User-Item matrix
- Columns sum to 1 representing probability from item \rightarrow user



$$P_{UI} = \begin{bmatrix} 0.32 & 0.22 & 0 & 0.06 & 0.13 & 0 & 0.17 & 0 & 0 & 0.09 \\ 0.39 & 0 & 0.12 & 0 & 0.25 & 0.21 & 0 & 0 & 0.04 & 0 \\ 0 & 0.24 & 0.19 & 0.12 & 0 & 0.14 & 0.09 & 0.05 & 0 & 0.16 \\ 0.18 & 0.24 & 0 & 0 & 0.15 & 0 & 0.18 & 0 & 0.08 & 0.16 \\ 0.29 & 0.16 & 0.14 & 0.09 & 0 & 0.19 & 0 & 0.07 & 0.05 & 0 \end{bmatrix}$$

$$P_{UI}(j, i) = \frac{w'_{ij}}{\sum_k w'_{kj}}$$

$$P_{UI} = \begin{bmatrix} 0.25 & 0.21 & 0 & 0.20 & 0.22 & 0 & 0.33 & 0 & 0 & 0.18 \\ 0.29 & 0 & 0.21 & 0 & 0.41 & 0.34 & 0 & 0 & 0.17 & 0 \\ 0 & 0.28 & 0.45 & 0.46 & 0 & 0.27 & 0.21 & 0.41 & 0 & 0.38 \\ 0.19 & 0.32 & 0 & 0 & 0.37 & 0 & 0.46 & 0 & 0.53 & 0.44 \\ 0.27 & 0.19 & 0.34 & 0.34 & 0 & 0.39 & 0 & 0.59 & 0.30 & 0 \end{bmatrix}$$

$$P_{UI}(i, j) = \frac{w'_{ij}}{\sum_k w'_{ik}}$$

Personalized PageRank - Components

- **Item Popularity**

- Log transformation followed by normalization
- Each item is assigned a popularity score
- Provides a fallback for users with short history

- **Personalization Vector**

- Initialized based on user's seed items
- Adjusted for users with varying interaction histories

$$W' = \begin{bmatrix} 0.87 & 0.60 & 0 & 0.17 & 0.35 & 0 & 0.47 & 0 & 0 & 0.24 \\ 1.00 & 0 & 0.30 & 0 & 0.65 & 0.54 & 0 & 0 & 0.10 & 0 \\ 0 & 0.78 & 0.63 & 0.40 & 0 & 0.44 & 0.30 & 0.17 & 0 & 0.51 \\ 0.65 & 0.90 & 0 & 0 & 0.57 & 0 & 0.67 & 0 & 0.30 & 0.60 \\ 0.94 & 0.54 & 0.47 & 0.30 & 0 & 0.63 & 0 & 0.24 & 0.17 & 0 \end{bmatrix}$$

$$\text{pop}_j = \sum_i w'_{ij}$$

$$\text{pop}'_j = \log(1 + \text{pop}_j)$$

$$\text{pop}''_j = \frac{\text{pop}'_j}{\sum_k \text{pop}'_k}$$

Game	Raw Sum	Log(1+sum)	Normalized
CSGO	3.46	1.50	0.15
Dota 2	2.82	1.34	0.13
Elden Ring	1.40	0.88	0.09
Minecraft	0.87	0.63	0.06
GTA V	1.57	0.95	0.09
RDR2	1.61	0.96	0.10
Witcher 3	1.44	0.89	0.09
Portal 2	0.41	0.34	0.03
Half-Life 2	0.57	0.45	0.04
Cyberpunk 2077	1.35	0.85	0.08

Personalized PageRank – Personalization Vector

- **Consider a new user**
 - Elden Ring - 45 hours
 - The Witcher 3 - 30 hours
 - GTA V - 15 hours
- Playtime Processing
- Initial Personalization Vector

$$t = \begin{bmatrix} 45 \\ 30 \\ 15 \end{bmatrix} \quad t_{\log} = \log(1 + t) = \begin{bmatrix} 3.83 \\ 3.43 \\ 2.77 \end{bmatrix}$$

$$v_j = \begin{cases} t_{\text{norm},j} & \text{if game } j \text{ is in history} \\ 0 & \text{otherwise} \end{cases}$$

$$v = [0 \quad 0 \quad 0.38 \quad 0 \quad 0.28 \quad 0 \quad 0.34 \quad 0 \quad 0 \quad 0]^T$$


Personalized PageRank - Components

- **Power Iteration Process:**

- Iteratively propagate personalization vectors
- $R_{next} = \alpha(R \cdot P_{IU} \cdot P_{UI}) + (1 - \alpha)v_p$
- Stop when convergence criteria hits or after **num_iterations**

- **Adaptive Popularity Weights:**

- Adjust based on user history length
- Cold-Start users \rightarrow Rely entirely on item popularity
- Short history \rightarrow Higher weight on popularity
- Long History \rightarrow Base popularity gets maintained
- For history of 3 games:
 - $\omega(3) = \min(0.5, \omega_b + (5 - 3) \cdot 0.1) = 0.4$



$$\omega(h) = \begin{cases} 0.9 & \text{if } h = 0 \text{ (no history)} \\ \min(0.5, \omega_b + (5 - h) \times 0.1) & \text{if } 1 \leq h \leq 5 \\ \omega_b & \text{if } h > 5 \end{cases}$$

Personalized PageRank - Recommendations

- **Final Scores:**

- Weighted combination of PPR and item popularity
- $score_j = (1 - \omega(h))R_{final,j} + \omega(h) \cdot pop''_j$
- $score_j = (1 - 0.4)R_{final,j} + 0.4 \cdot pop''_j$

- **Recommendations:**

- Exclude known interactions, zero out
- Select top-n items

$score_{final} =$

CSGO: 0.16
Dota 2: 0.16
Elden Ring: 0.24
Minecraft: 0.07
GTA V: 0.20
RDR2: 0.17
Witcher 3: 0.23
Portal 2: 0.08
Half-Life 2: 0.09
Cyberpunk: 0.15

MultiAlphaPPR

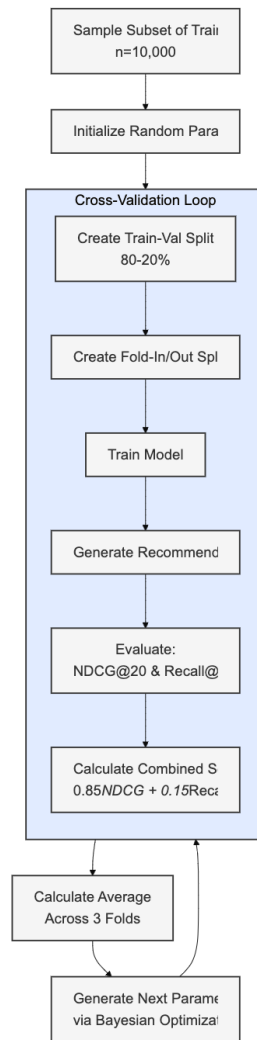
- Extends PPR by utilizing multiple damping factors
- Combines multiple weighted PPR scores
- Additional Parameters:
 - **Alphas:**
 - List of damping factors
 - **Alpha Weights:**
 - Corresponding weights for each alpha, summing up to 1
- **MultipleRandom Walks:** run PPR independently for each alpha
- **Weighted Combination:** aggregate PPR scores using alpha weights
- **Final Scoring:** Combine aggregated PPR with item popularity

TwoStagePPR

- Enhances PPR with two-phase approach
- Combines broad exploration and focused exploitation
- Additional Parameters:
 - **Stage 1 Alpha (α_1):**
 - Higher damping factor for broad exploration (e.g., 0.7)
 - **Stage 2 Alpha (α_2):**
 - Lower damping factor for focused exploitation (e.g., 0.3)
 - **Stage 1 Top-K Items:**
 - Number of top items to consider from first stage
- Initialize personalization with top-K items
 - Run PPR with α_2 to refined recommendations
- Final scoring averages PPR from both stages

Optimization

- **Bayesian Optimization + Sampling:**
 - Efficient exploration of hyperparameter space
- Evaluation using weighted combination of NDCG@20 & Recall@20
 - $0.85\text{NDCG@20} + 0.15\text{Recall@20}$
- **Process:**
 - Sample dataset while maintaining distributional characteristics
 - Cross-Validation trials on different seeds:
 - Base seed calculates: sampling, cv and optimization seed
 - Model training & evaluation
 - Iterate to find optimal hyperparameters
- **Robustness:**
 - Different algorithms make use of same base seed
 - Same data split across trials




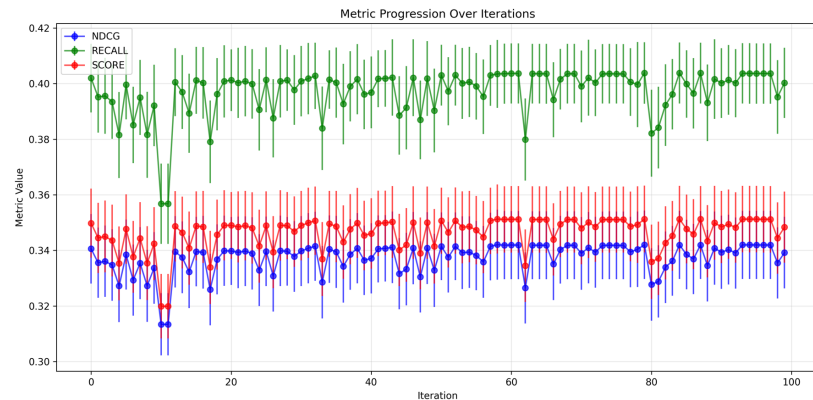
Personalized PageRank - Results

- **Optimized Parameters:**

- Alpha: 0.023
- Popularity Weight: 0.033
- num_iterations: 93
- Playtime processing: Log transformation followed by MinMaxScaler


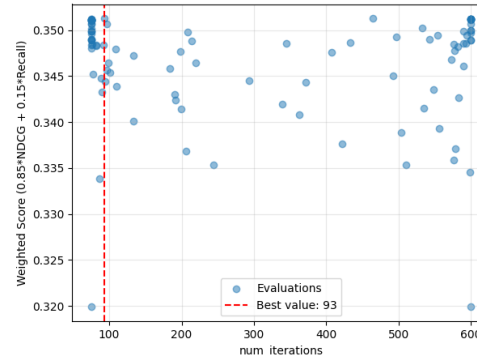

- **Performance:**

- Offline: NDCG@20 = 0.3449, Recall@20 = 0.4172
- Offline: NDCG@20 = 0.3183, Recall@20 = 0.3790
- Execution Time: 280s (CPU), 30-45s (GPU)
- Outperform all baselines except Item-KNN
- Comparable to baseline execution times



Personalized PageRank - Results

- **Low Alpha Implications:**
 - Minimal reliance on random walk
 - Greater influence of personalization vector
 - Potential underutilization of graph structure
- **Popularity Weight Near Zero:**
 - Emphasizes personalized scores over global popularity



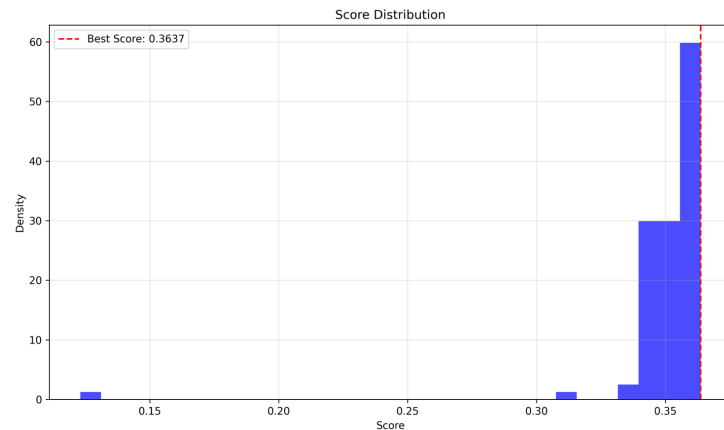

TwoPhasePPR - Results

- **Optimized Parameters:**

- Alpha1: 0.041
- Alpha2: 0.0
- Stage1_K: 148
- Popularity Weight: 0.0
- num_iterations: 271

- **Performance:**

- Offline: NDCG@20 = 0.3453, Recall@20 = 0.4175
- Offline: NDCG@20 = 0.3181, Recall@20 = 0.3790
- Execution Time: 350 (CPU), 40-55s (GPU)
- Gives almost the exact same results
- Slightly longer than PPR due to additional phase (that actually has no influence...)



TwoPhasePPR - Results

- **Alpha2 at zero:**



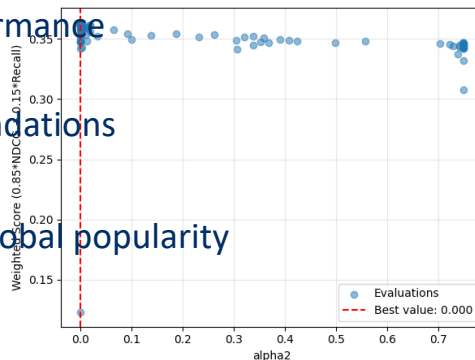
- Second phase no influence, effectively reducing to single-phase PPR
- No contribution → redundant
- No additional performance gains observed from 2nd phase
- Extra complexity != extra performance

- **Stage1_K Impact:**


- Moderate variation, no significant performance difference across different K values
- Suggests minimal impact on recommendations

- **Popularity Weight near zero:**

- Emphasizes personalized scores over global popularity



Dicussion – Baseline Comparison



Discussion – Alpha Insights

- **Effect of Low Alpha:**
 - Dominance of personalization vector over graph exploration
 - Minimal influence of the link structure
 - Underutilization of graph data
- **Alpha = 0:**
 - Comparable to Popularity baseline
 - No graph propagation leading to popularity-based suggestions
- **Small Alpha vs. Zero Alpha:**
 - Slight influence of graph structure can still enhance recommendations
 - Recommendations degrade to user history or popularity (no graph propagation), leading to popularity-based suggestions
- **MultiAlpha:**
 - Low alphas dominate, rendering multiple alphas ineffective
 - Leads to redundant computations

Dicussion - Why GNNs where considered

- **Strengths:**

- Capable of capturing complex user-item relationships through multi-hop message passing
- Enhanced representation learning by aggregating informations from neighbours

- **Challenges:**

- GNNs like LightGCN rely on pre-learned user embeddings from user data
- Unable to generate embeddings for new users without retraining
- Significant computational resources needed for large-scale graphs due to multi-hop message passing
- Modified GNNs did not outperform simpler methods like PPR and Item-KNN in small-scale testing

- ➔ Discontinue GNN research

Future Work & Improvements

- **Modified GNN with Item Co-occurrence:**
 - Create co-occurrence graph based on shared user interactions
 - Capture higher-order relations
 - Achieved similar performance to PPR and Item-KNN in small-scale testing
 - Remains highly resource-intensive (4-5 hours of training with graph reduction techniques)
- **Hybridizing PPR with Content-Based Features:**
 - Merge PPR with content-based filtering to leverage both interaction data and rich item attributes
- **Promoting Diversity & Fairness:**
 - Varied set of recommendations to avoid over-reliance on popular items